Understanding Databricks/Spark Performance Tuning

Lesson 02: Easiest Fix – Optimizing the Hardware

by Bryan Cafferky from my YouTube channel

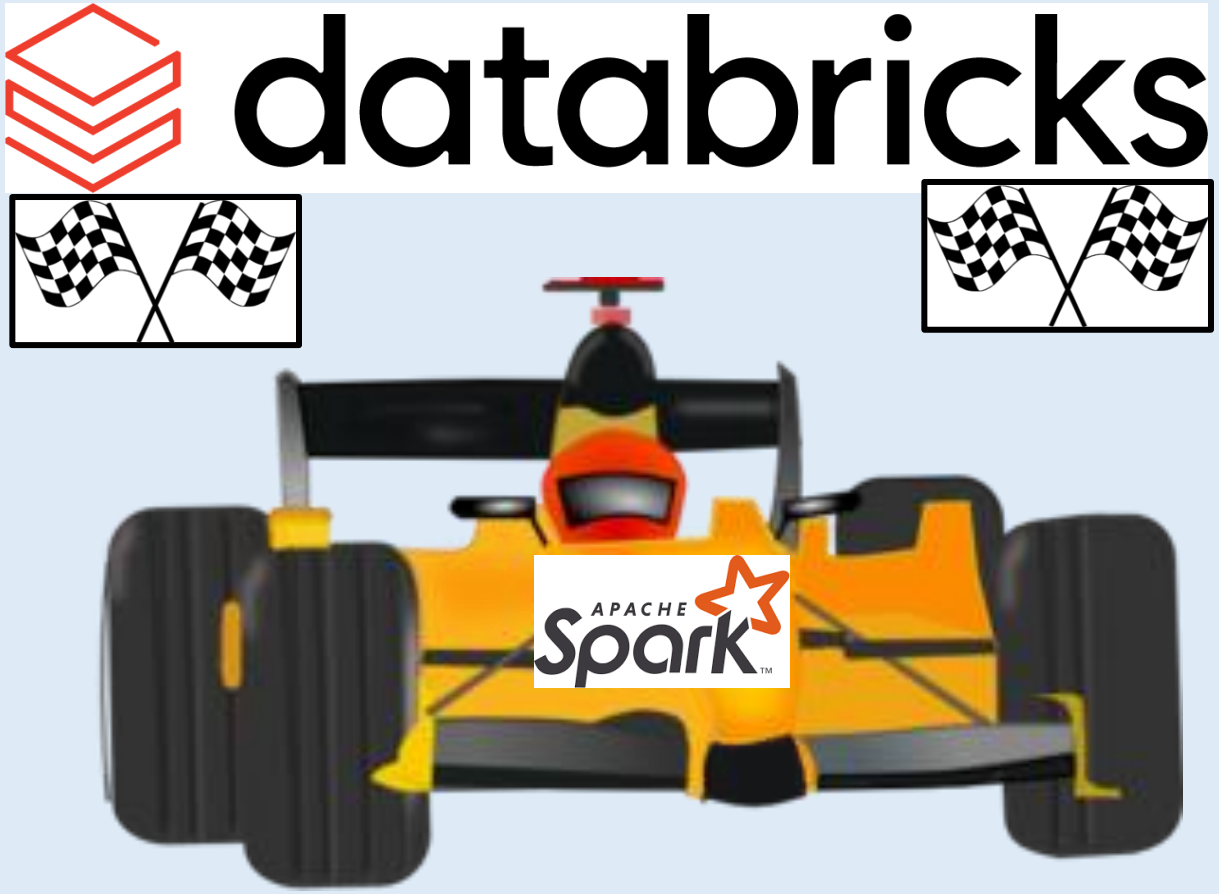

Where We're Going?

- Why Compute Resources Matter
- Databricks Workspace – Standard or Premium
- Databricks/Apache Spark Cluster Architecture
- Hardware Under the Cluster Architecture
- Optimizing the Cluster Configuration Step By Step
- Shuffles & Spills
- Wrap Up

Why Compute Resources Matter

- Never Talked About in Depth
- Cost vs. Performance Trade Off
- Single Easiest Thing You Can Change to Improve Performance
- There are Interdependencies Between Performance Optimization & the Compute Selection/Options

Databricks Workspace

Databricks Workspace

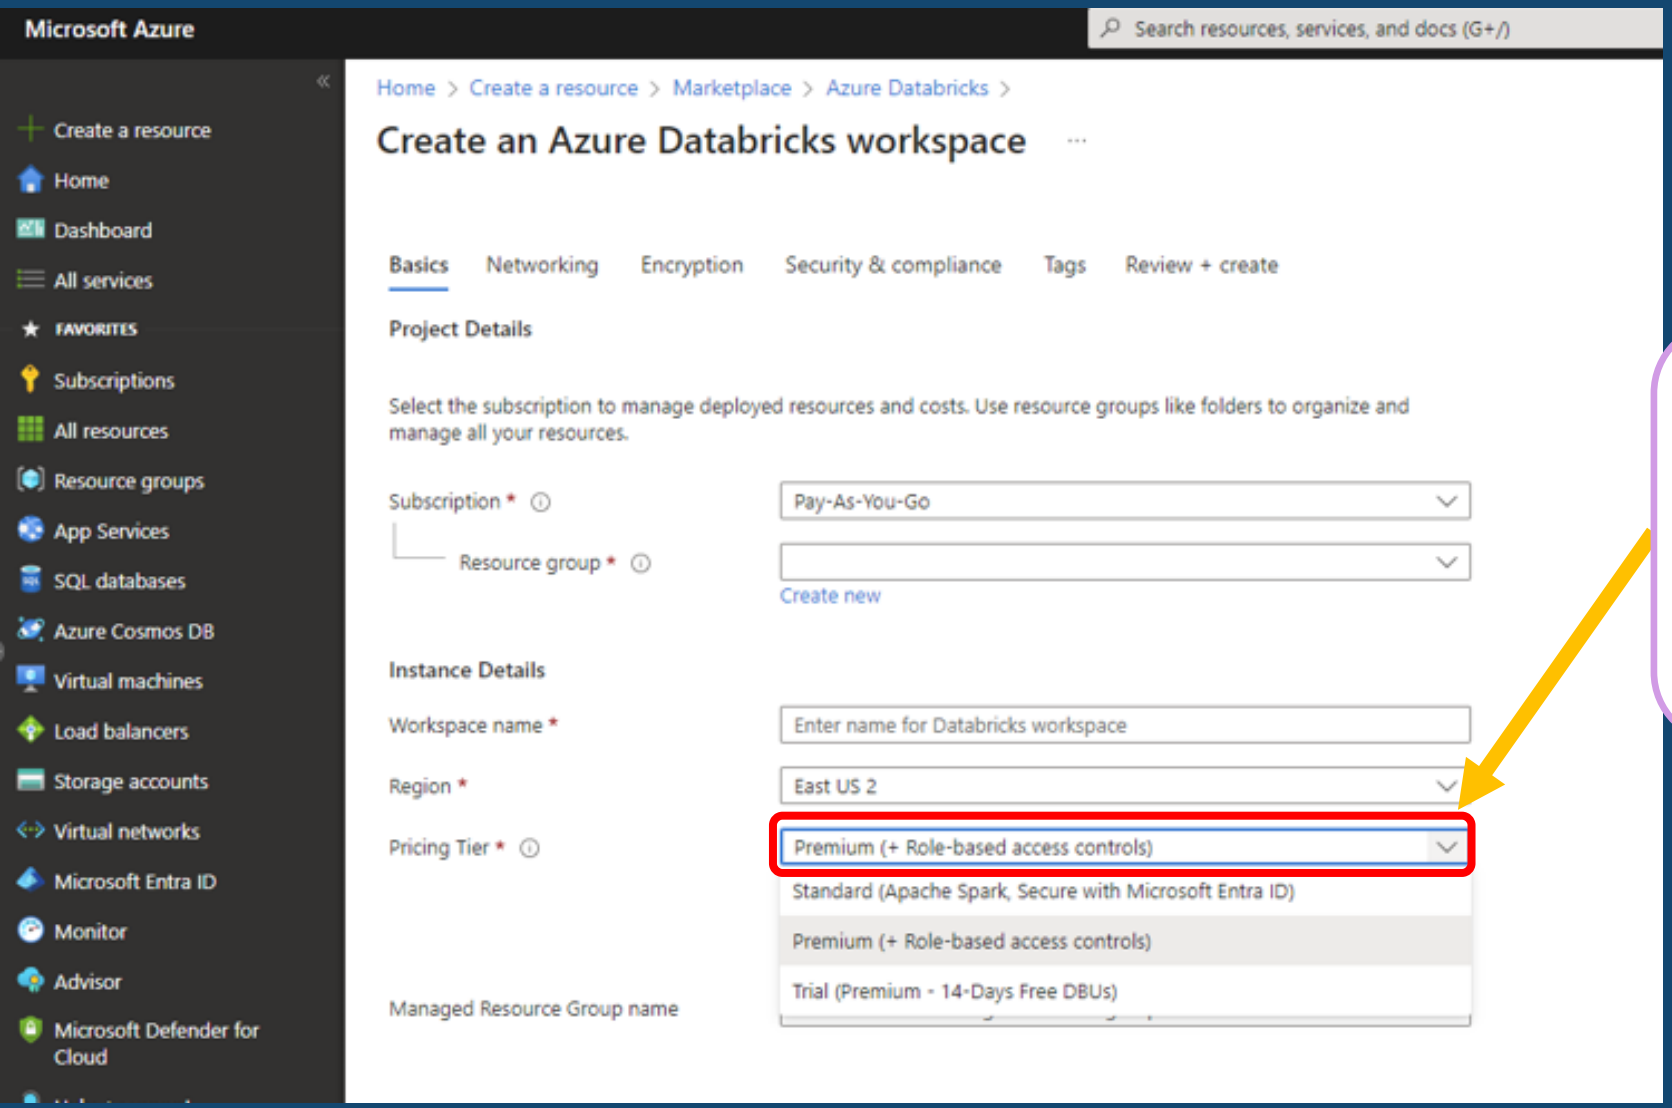

Premium vs. Standard

- Unity Catalog Requires Premium
- Delta Live Tables Requires Premium
- Photon Does NOT Require Premium

Unity Catalog

Before you begin

Before you begin the tasks described in this article, you should familiarize yourself with the basic Unity Catalog concepts, including metastores, admin roles, and managed storage. See What is Unity Catalog?.

You should also confirm that you meet the following requirements:

- An Azure Databricks workspace on the Premium plan.
- The following roles and privileges, which depend on the status of your workspace:
  - Workspace admin: If your workspace was enabled for Unity Catalog automatically when it was created, you must be a workspace admin to complete the required tasks.
  - Account admin: If your workspace is not already enabled for Unity Catalog, an account admin must attach the workspace to the metastore.

    If there is no Unity Catalog metastore in the same region as the workspace, an account admin must also create the Unity Catalog metastore.

    Instructions for determining whether a metastore exists for your workspace region, along with instructions for creating a metastore, follow in this article.

See Admin privileges in Unity Catalog and Automatic enablement of Unity Catalog.

Premium Required for Unity Catalog

https://learn.microsoft.com/en-us/azure/databricks/data-governance/unity-catalog/get-started#before-you-begin

Delta Live Tables (DLT)

What is Delta Live Tables?

Article • 04/22/2024 • 3 contributors

Feedback

In this article

Delta Live Tables is a declarative framework for building reliable, maintainable, and testable data processing pipelines. You define the transformations to perform on your data and Delta Live Tables manages task orchestration, cluster management, monitoring, data quality, and error handling.

Premium Required for DLT

> **Note**
>
> Delta Live Tables requires the Premium plan. Contact your Databricks account team for more information.

https://learn.microsoft.com/en-us/azure/databricks/delta-live-tables/

Cluster Architecture

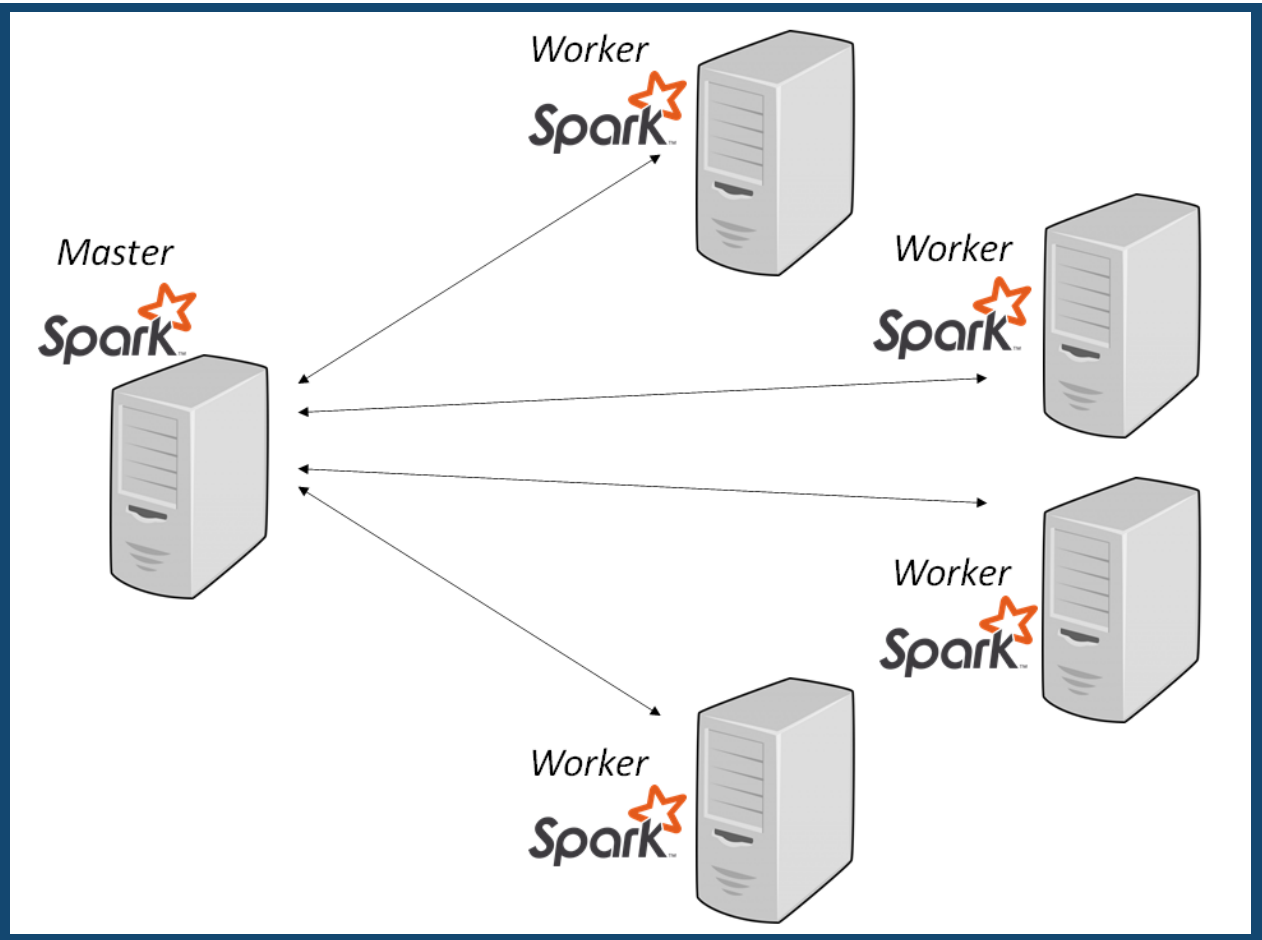

Apache Spark Core—Deep Dive—Proper Optimization Daniel Tomes Databricks

Databricks
103K subscribers

Subscribed

https://www.youtube.com/watch?v=daXEp4HmS-E&t=2111s

Cluster Architecture

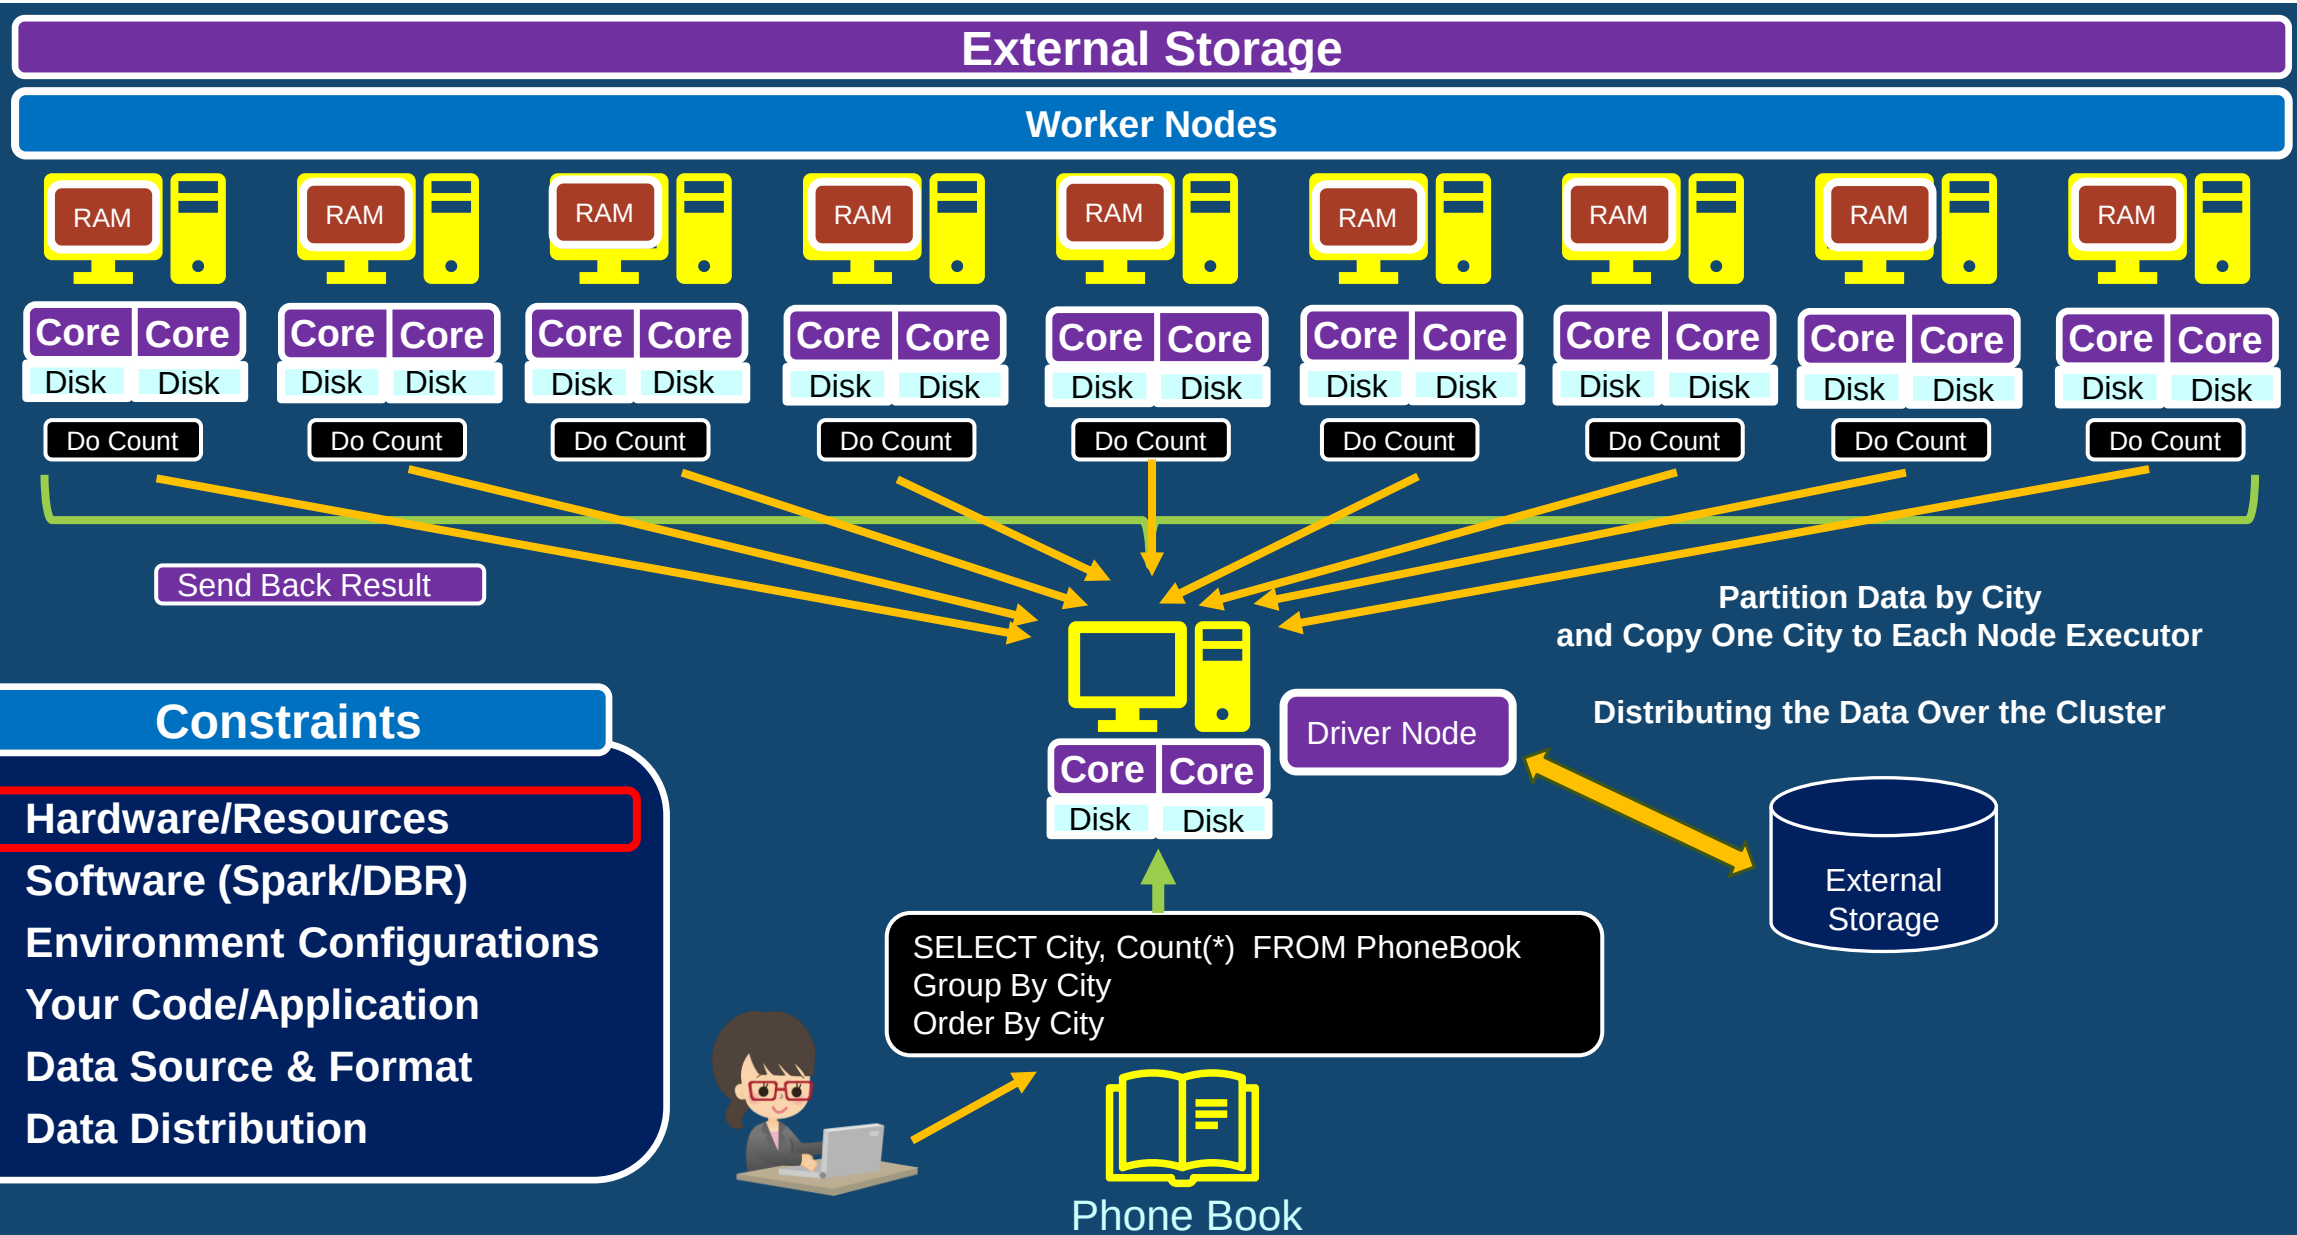

Constraints

- Hardware/Resources
- Software (Spark/DBR)
- Environment Configurations
- Your Code/Application
- Data Source & Format
- Data Distribution

Hardware Focus

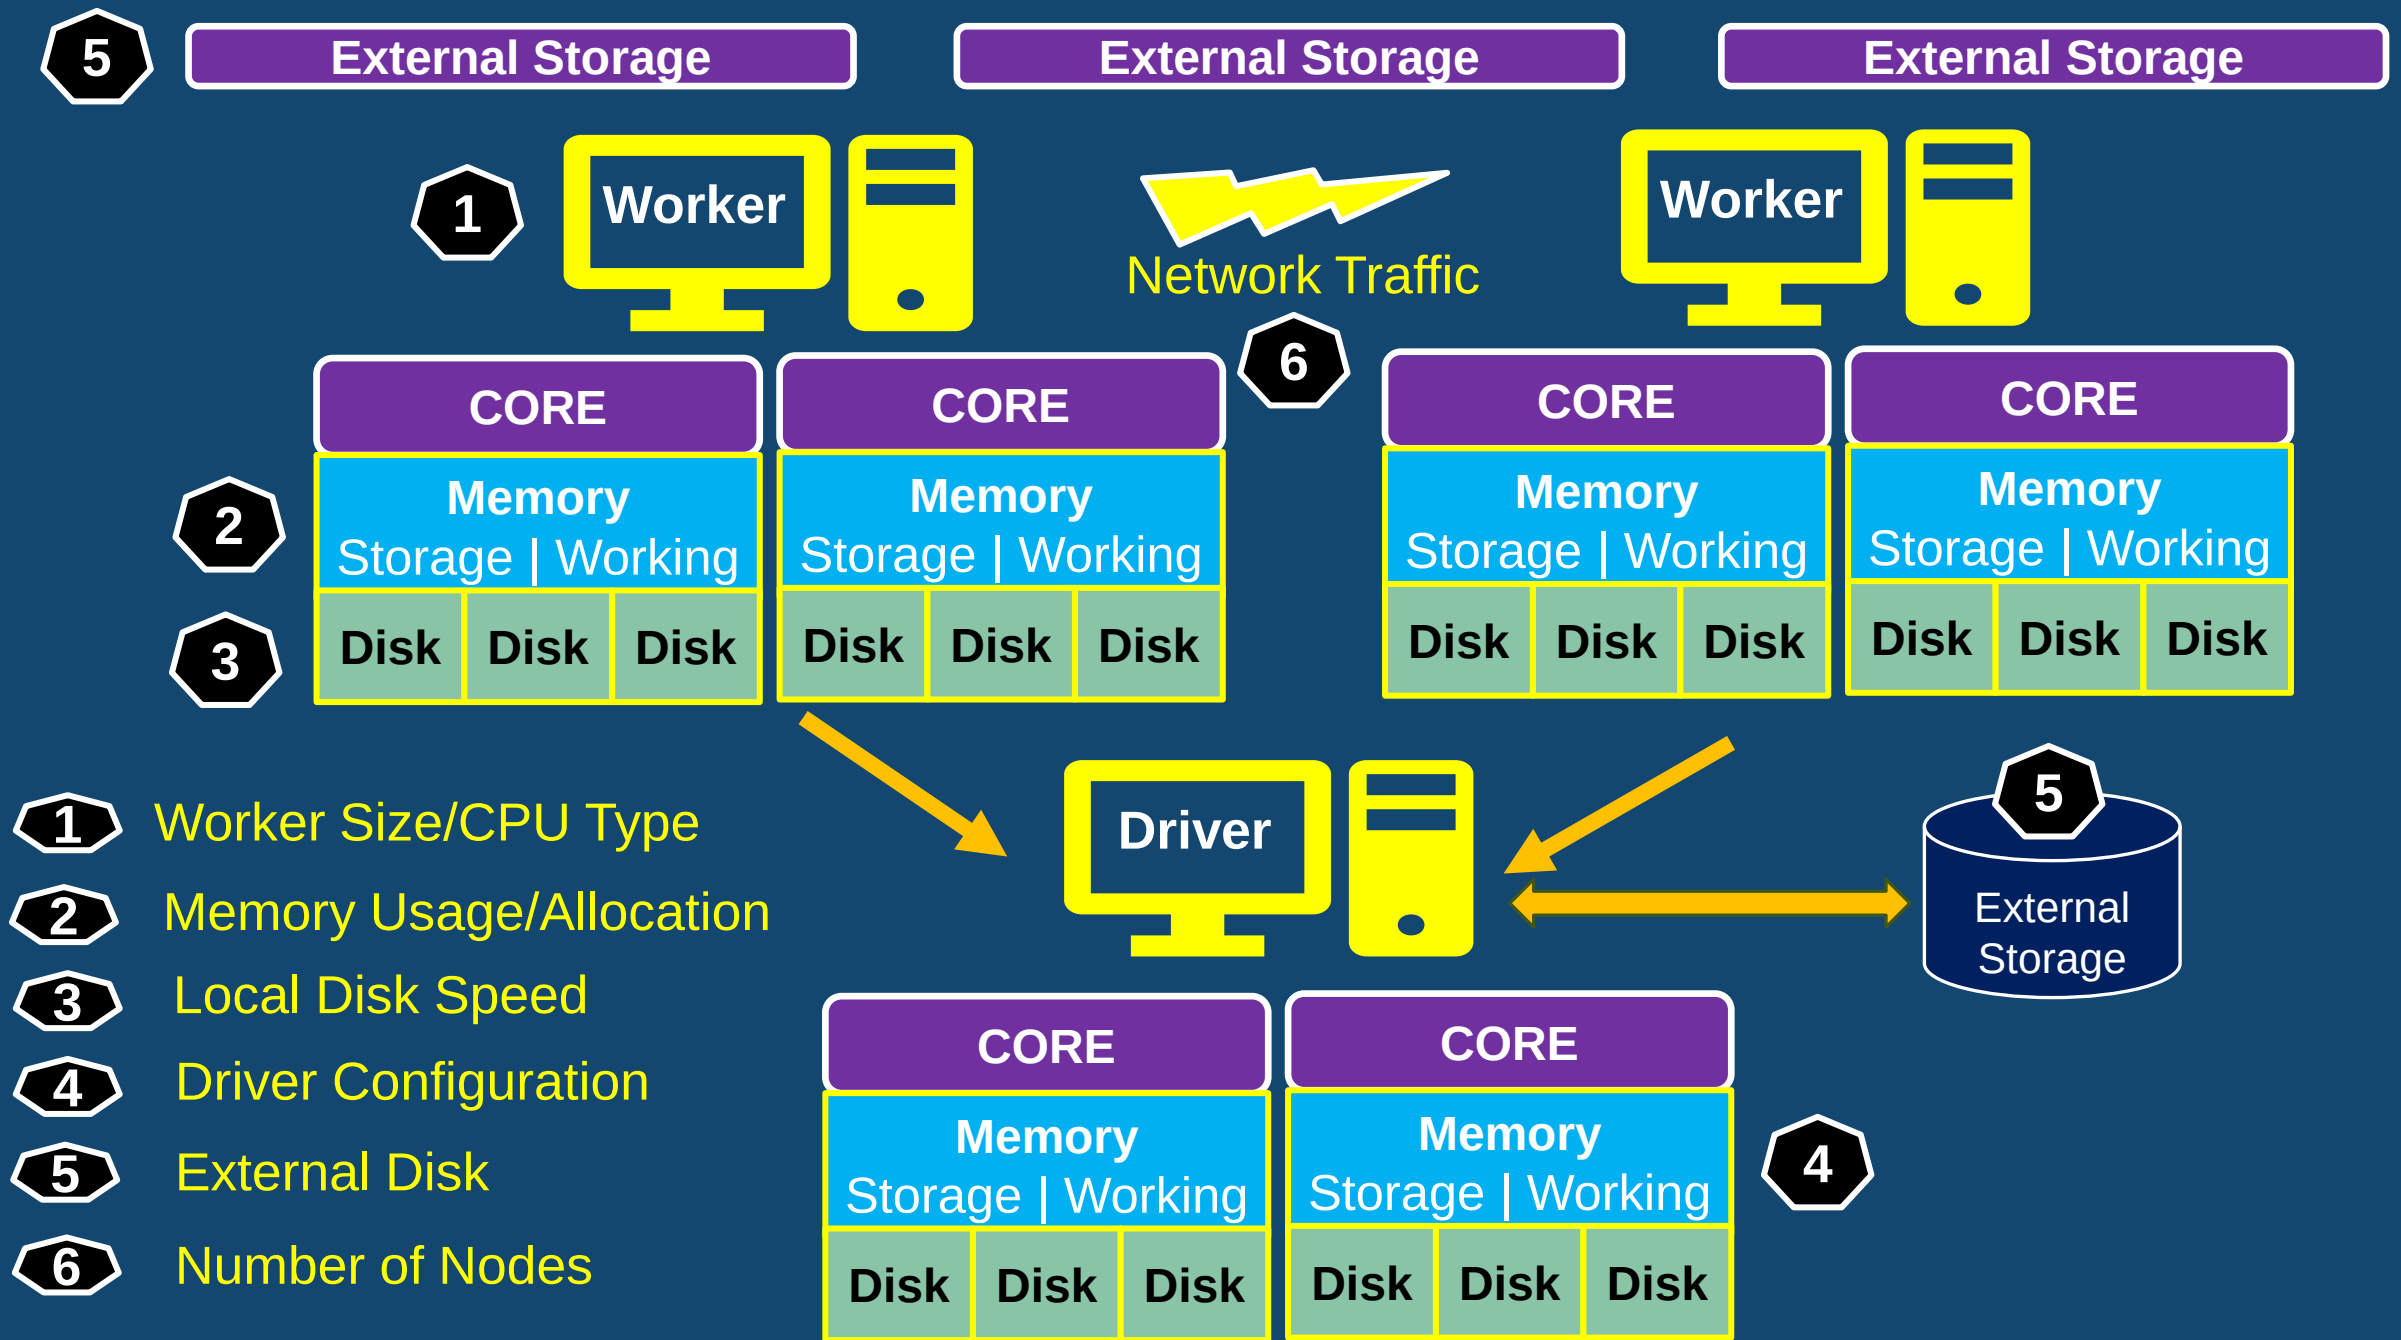

1. Worker Size/CPU Type
2. Memory Usage/Allocation
3. Local Disk Speed
4. Driver Configuration
5. External Disk
6. Number of Nodes

Hardware – CPU & Memory

- **CPU/GPU Type**
  - Workload Type
- **Number of Cores**
  - Core = Executor = Task = Partition
  - # of Cores = Degree of Parallelism
- **Memory**
  - Enough to Support Your Workload
  - Avoid Spills
  - Split Between Working & Storage Memory

Hardware – Input/Output

- **Network**
  - Latency (Local vs. Cross Region vs. On Prem)
  - Bandwidth
  - External Resource Latency (example: Cosmos DB vs Azure SQL)
- **Storage (Local, Remote (External))**
  - Type/Speed (HDD, SSD)
  - Locality
- **I/O**
  - Spills
  - Waiting for Disk Writes
- **Databricks Runtime Version (Effects Features)**

Optimizing the Cluster Configuration

Cluster – Configuration Settings

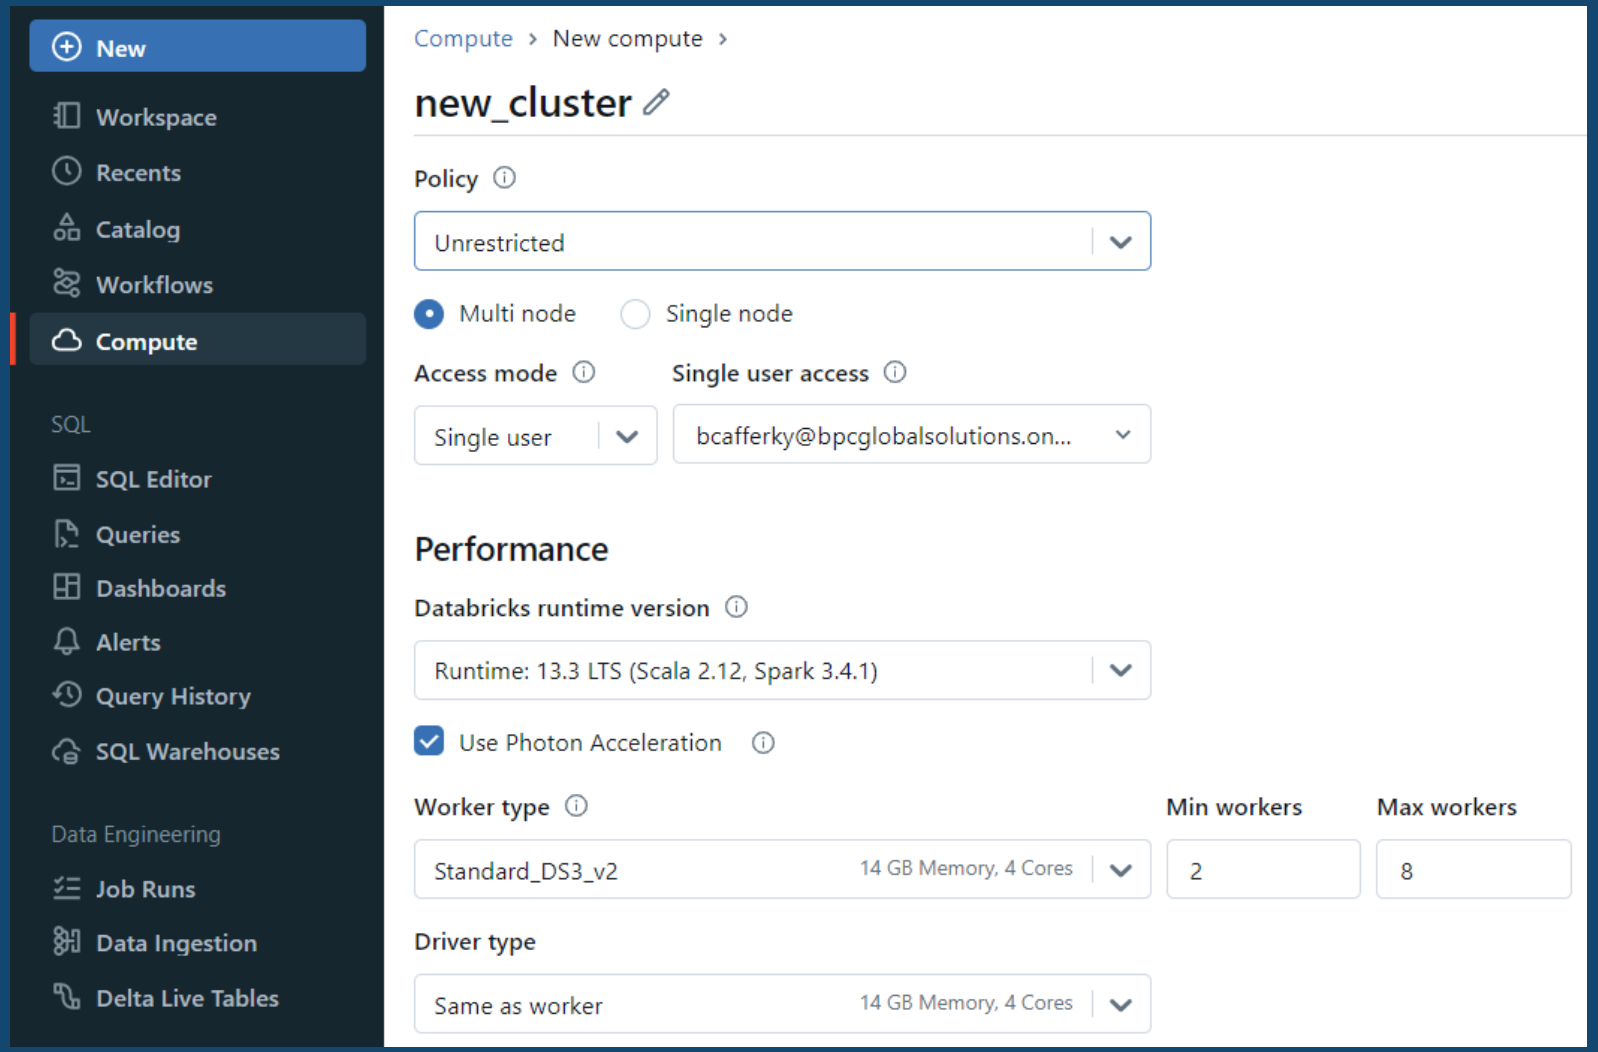

We'll Walk Through Filling Out This Screen.

https://docs.databricks.com/en/compute/configure.html

https://learn.microsoft.com/en-us/azure/databricks/compute/configure

Cluster – Configuration Best Practices

Use Case	Notes	Recommendation
Analysis 	Interactive Development	➢ Single Node with High Memory and Cores ➢ Likely require reading the same data repeatedly, so recommended node types are storage optimized with disk cache enabled.
Basic Batch ETL 	No Wide Transformations	Compute Optimized
Complex ETL 	Has Wide Transformations	Compute Optimized with Less Nodes
ML Training Experimentation/Dev 		Single Node Type with High Memory and Cores
ML Training Production 	Minimal Worker Nodes	Storage Optimized with Disk Caching Enabled or GPU (lacks disk caching)

https://learn.microsoft.com/en-us/azure/databricks/compute/cluster-config-best-practices

Cluster Options

Use Case	Notes	Recommendation
Spot Instances 	Saves money by uses available capacity. https://learn.microsoft.com/en-us/azure/virtual-machines/spot-vms	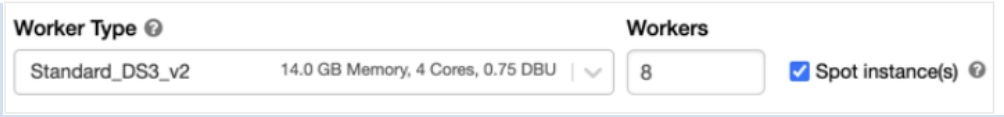
Serverless Compute 	Near instant Cluster. Unity Catalog must be enabled. https://learn.microsoft.com/en-us/azure/databricks/compute/serverless https://learn.microsoft.com/en-us/azure/databricks/release-notes/serverless#limitations https://www.databricks.com/trust/security-features/serverless-security	Like having a set of VMs on standby. When you need them, they are allocated to your work almost instantly.
Photon 	Vastly faster processing.	

Cluster – Access Mode

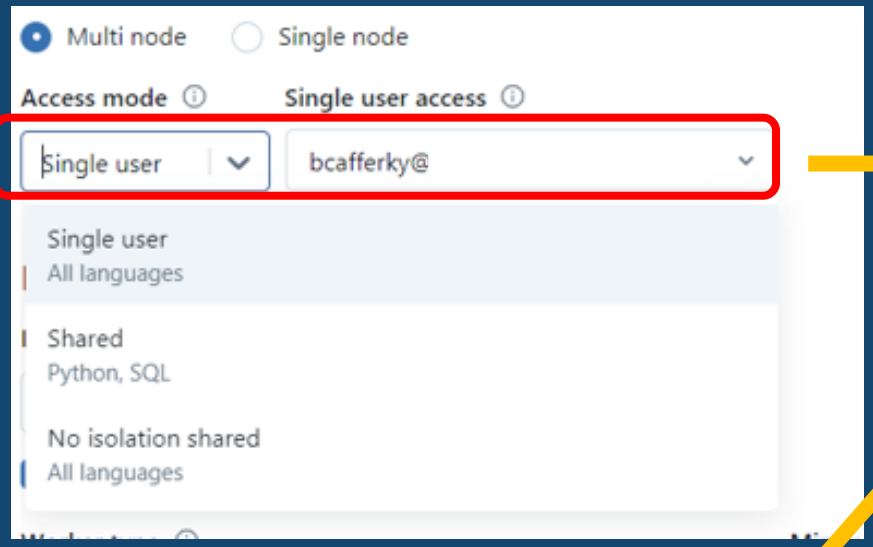

Single User for Credential Pass Through.

Credential Passthrough Automatically Passes Your Credentials Through to Backend Resources like ADLS.

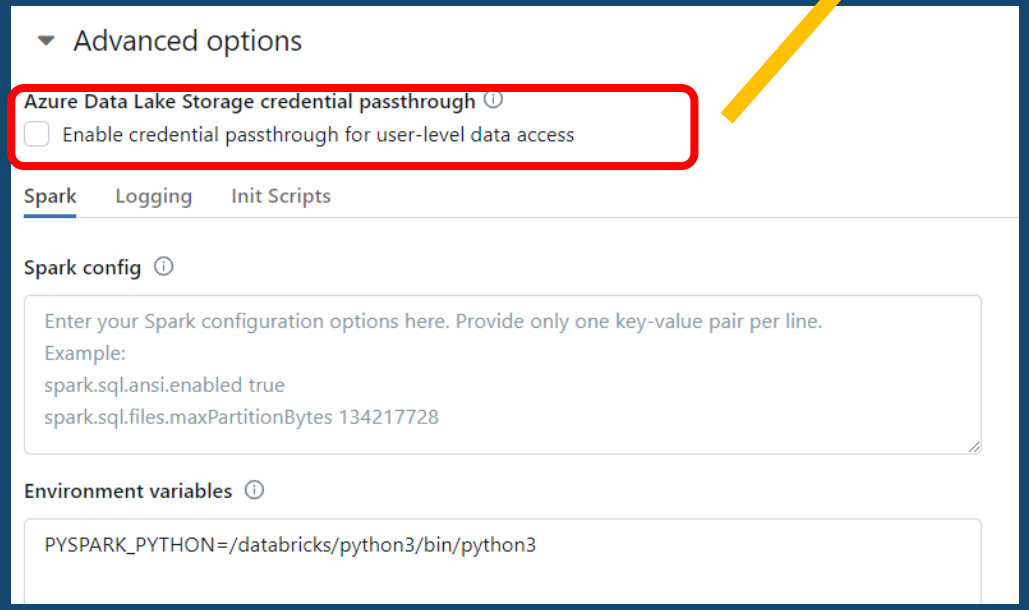

Cluster – Runtime and Node Types

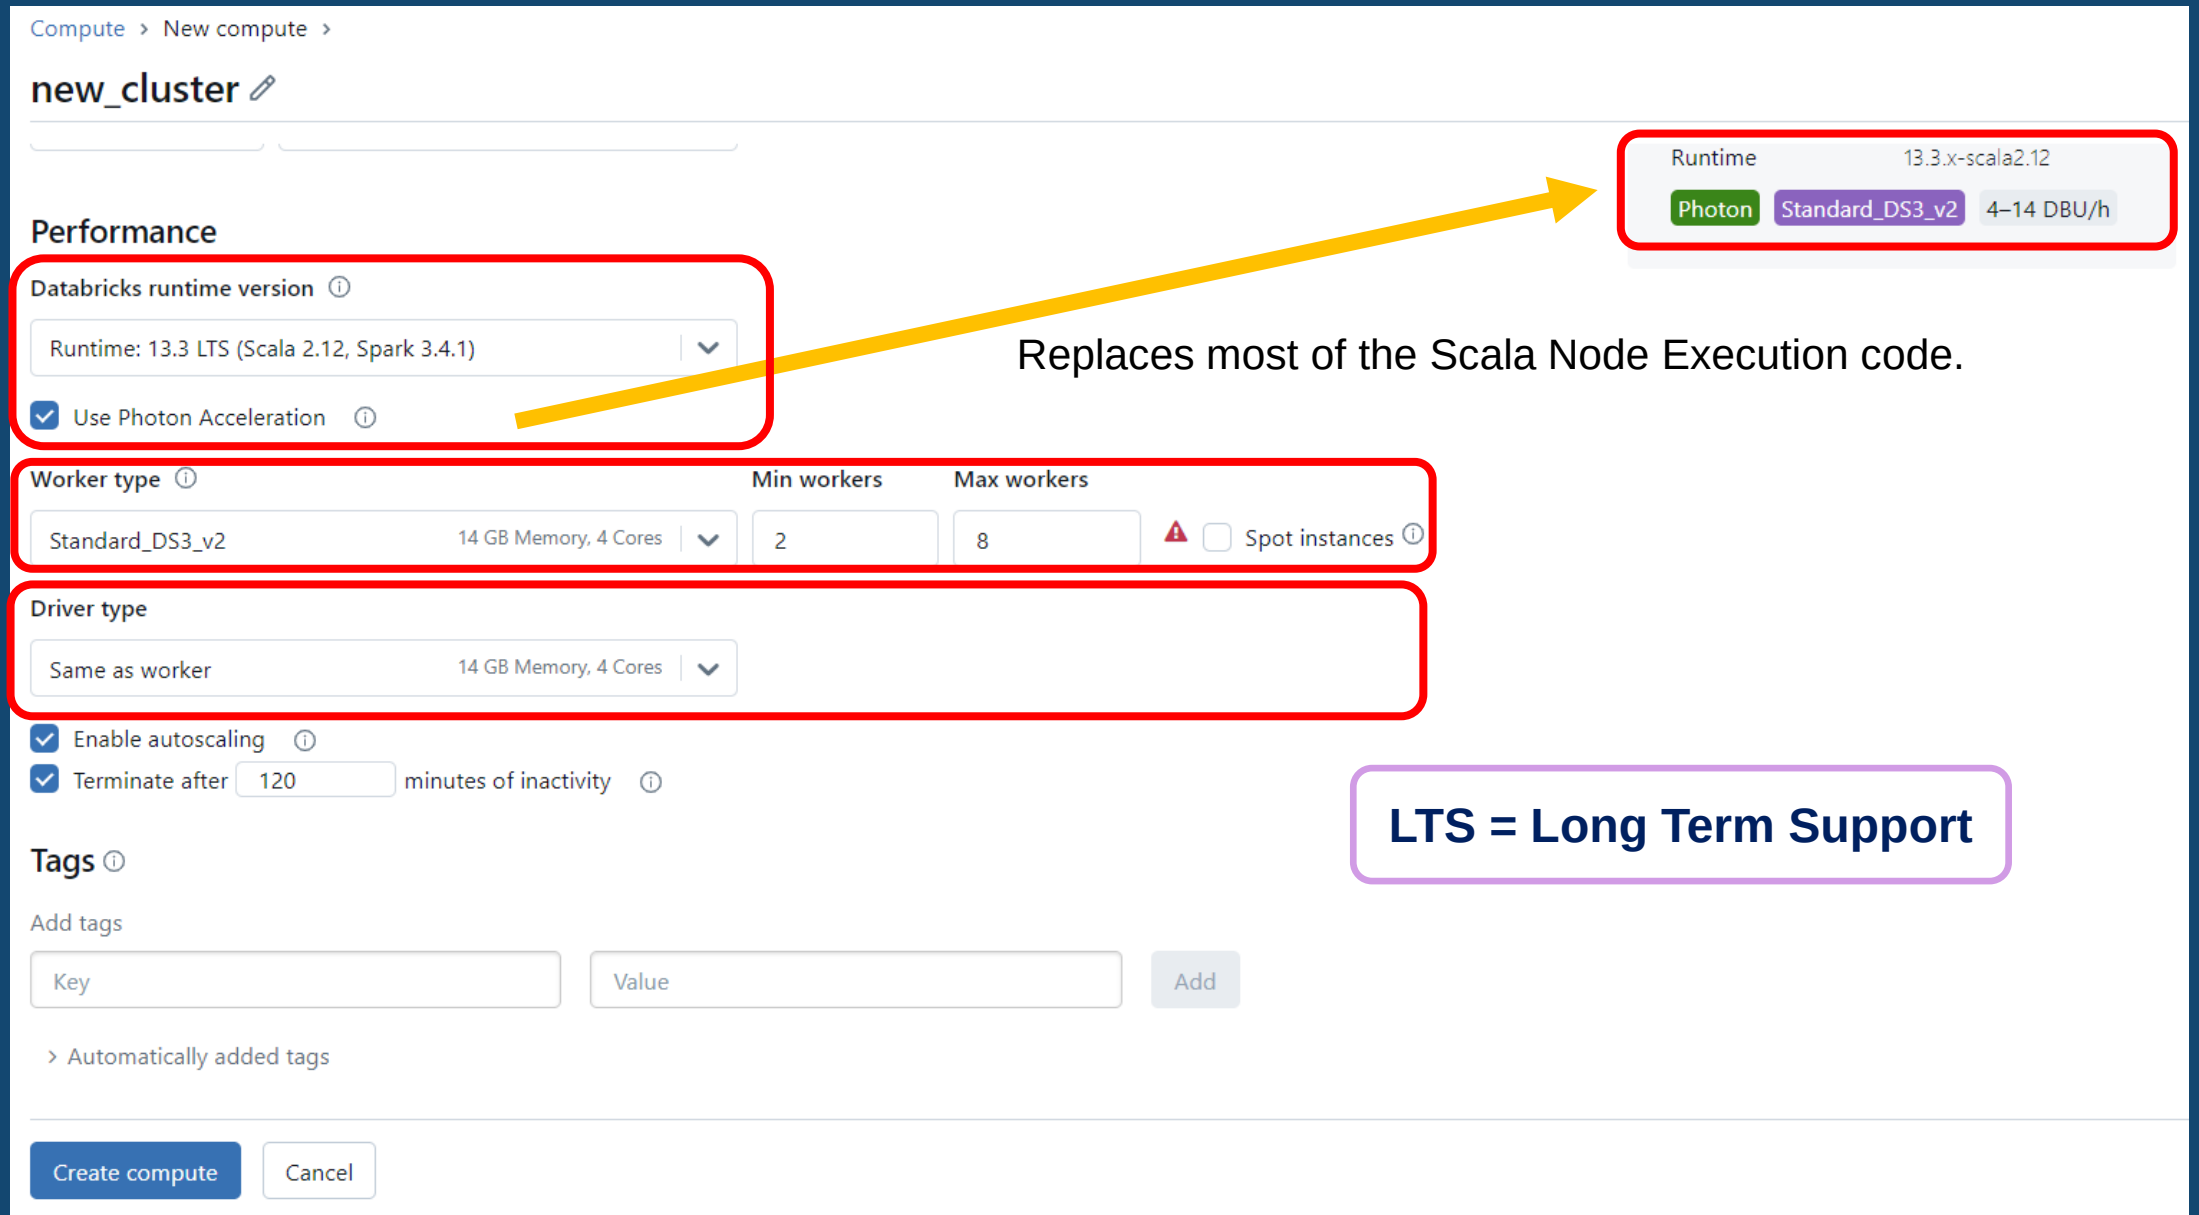

Photon – a New Execution Engine on Databricks

Get started with Photon

Photon is enabled by default on clusters running Databricks Runtime 9.1 LTS and above.

To manually disable or enable Photon on your cluster, select the **Use Photon Acceleration** checkbox when you create or edit the cluster.

If you create a cluster using the Clusters API, set `runtime_engine` to `PHOTON`.

Instance types

Photon supports a number of instance types on the driver and worker nodes. Photon instance types consume DBUs at a different rate than the same instance type running the non-Photon runtime. For more information about Photon instances and DBU consumption, see the Azure Databricks pricing page.

Photon is Only Available on Databricks

https://learn.microsoft.com/en-us/azure/databricks/compute/photon

Cluster – Delta cache accelerated

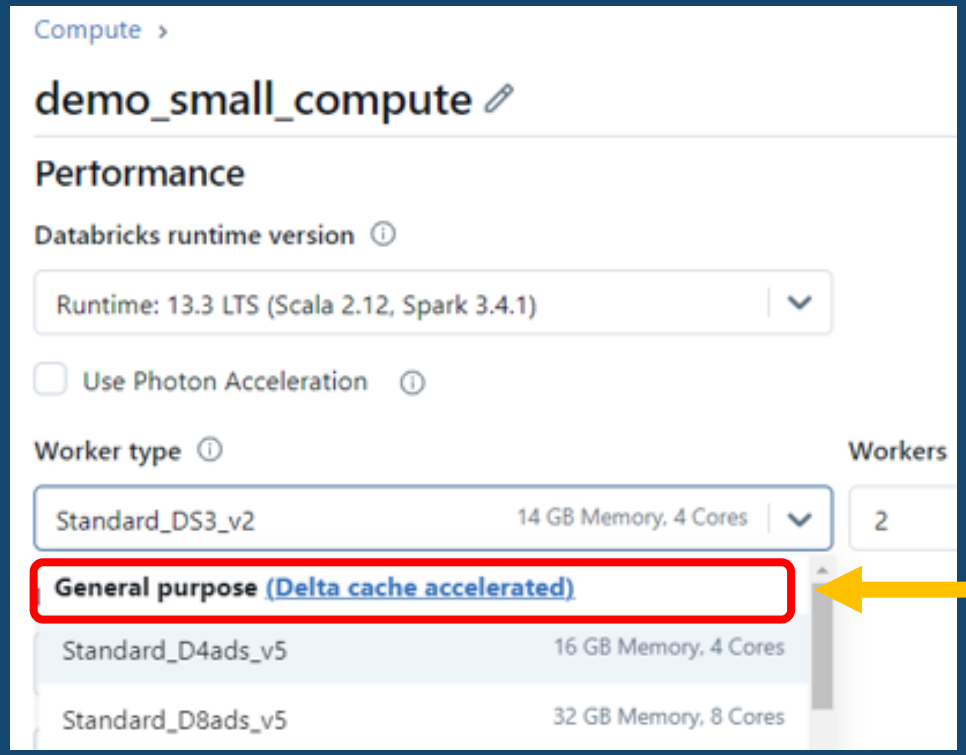

Improves Reading of Parquet and Delta Files.

Renamed to disk cache.

https://learn.microsoft.com/en-us/azure/databricks/optimizations/disk-cache

Cluster – Databricks Runtime

Standard

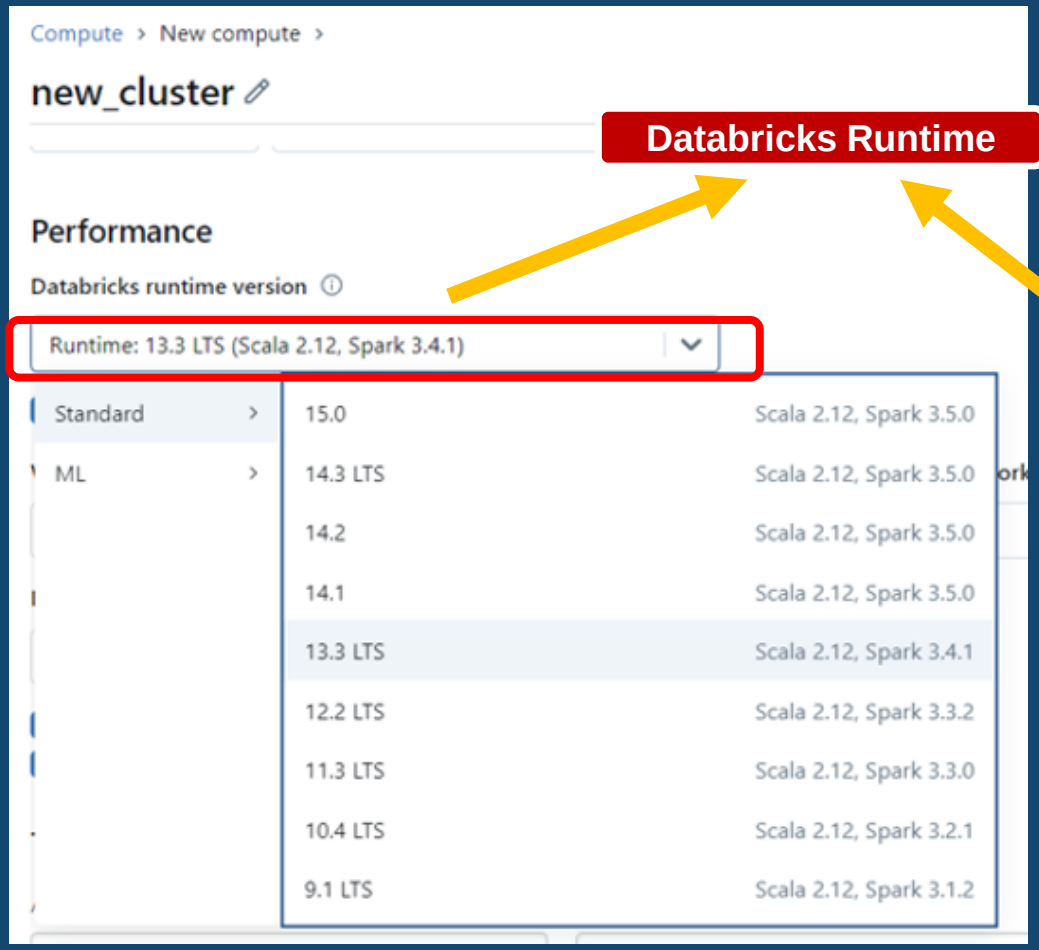

Databricks Runtime

Machine Learning

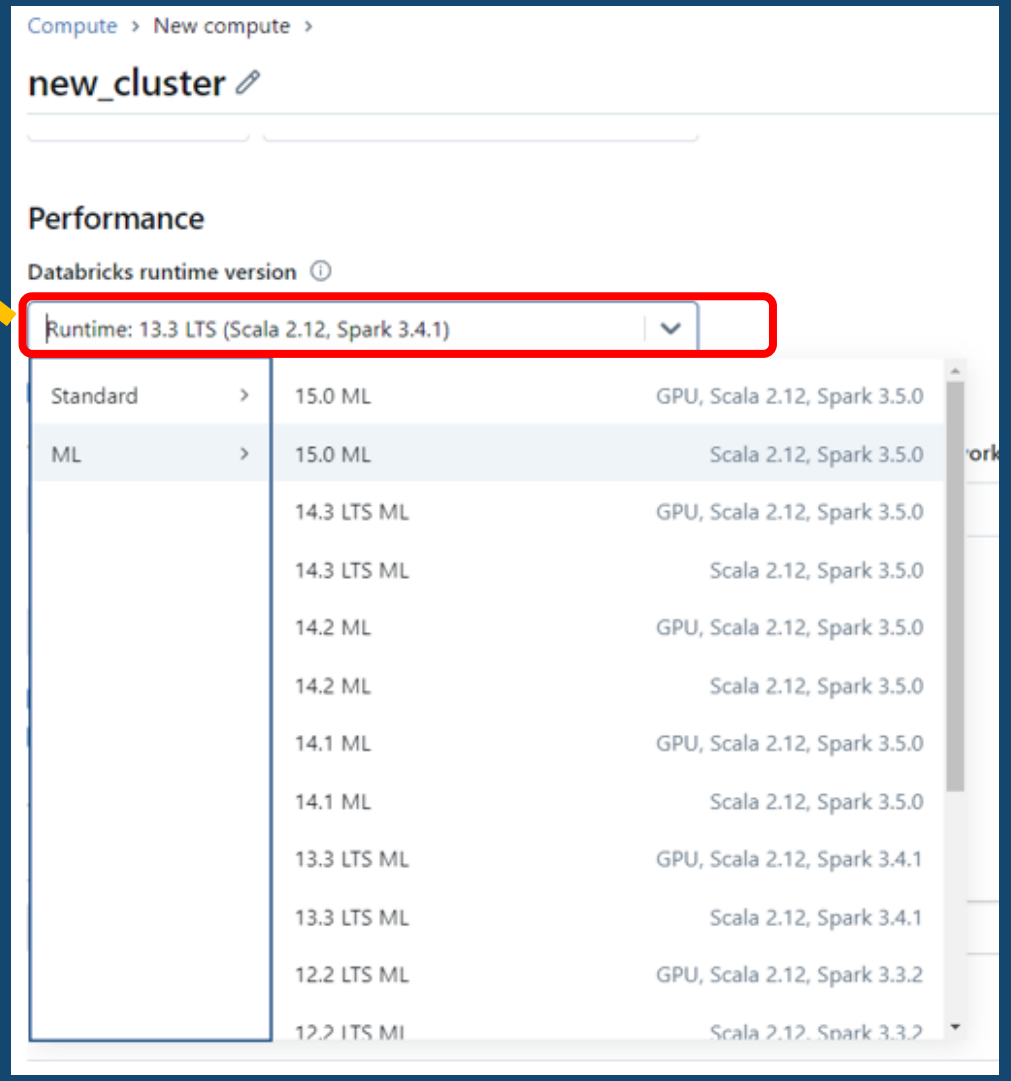

Cluster – Worker Type

Worker Type

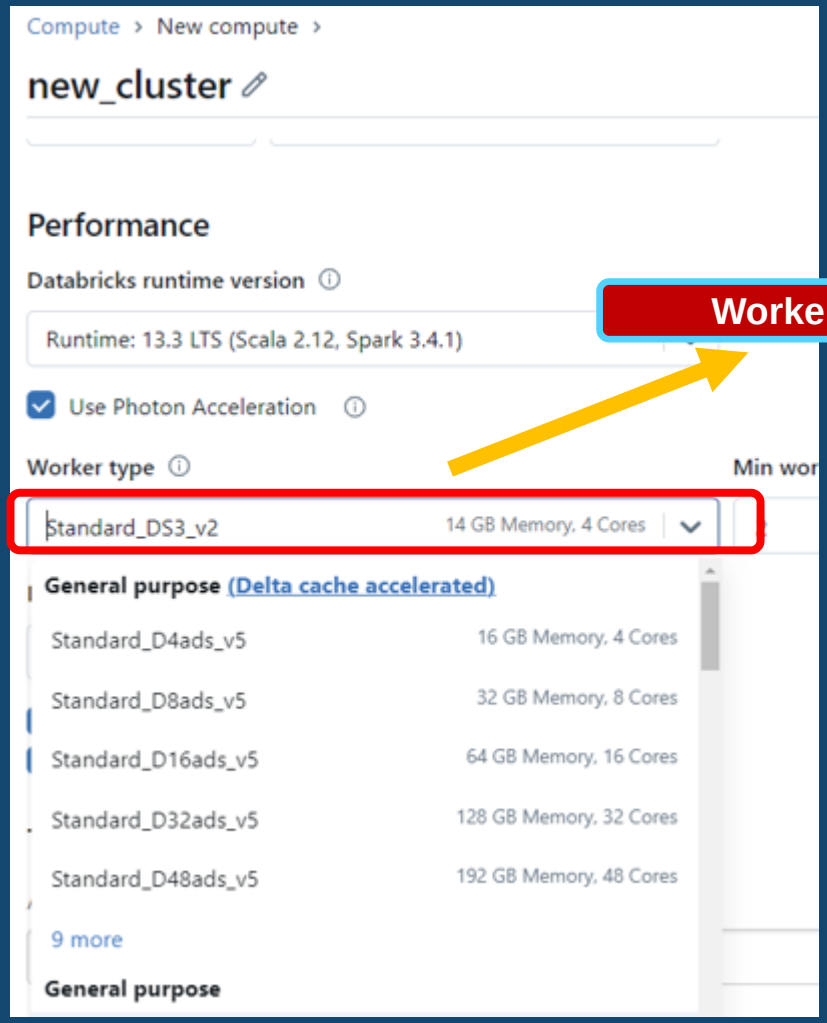

Worker Type

Larger Workers means less nodes required.

Cluster – Driver Type

Driver Type

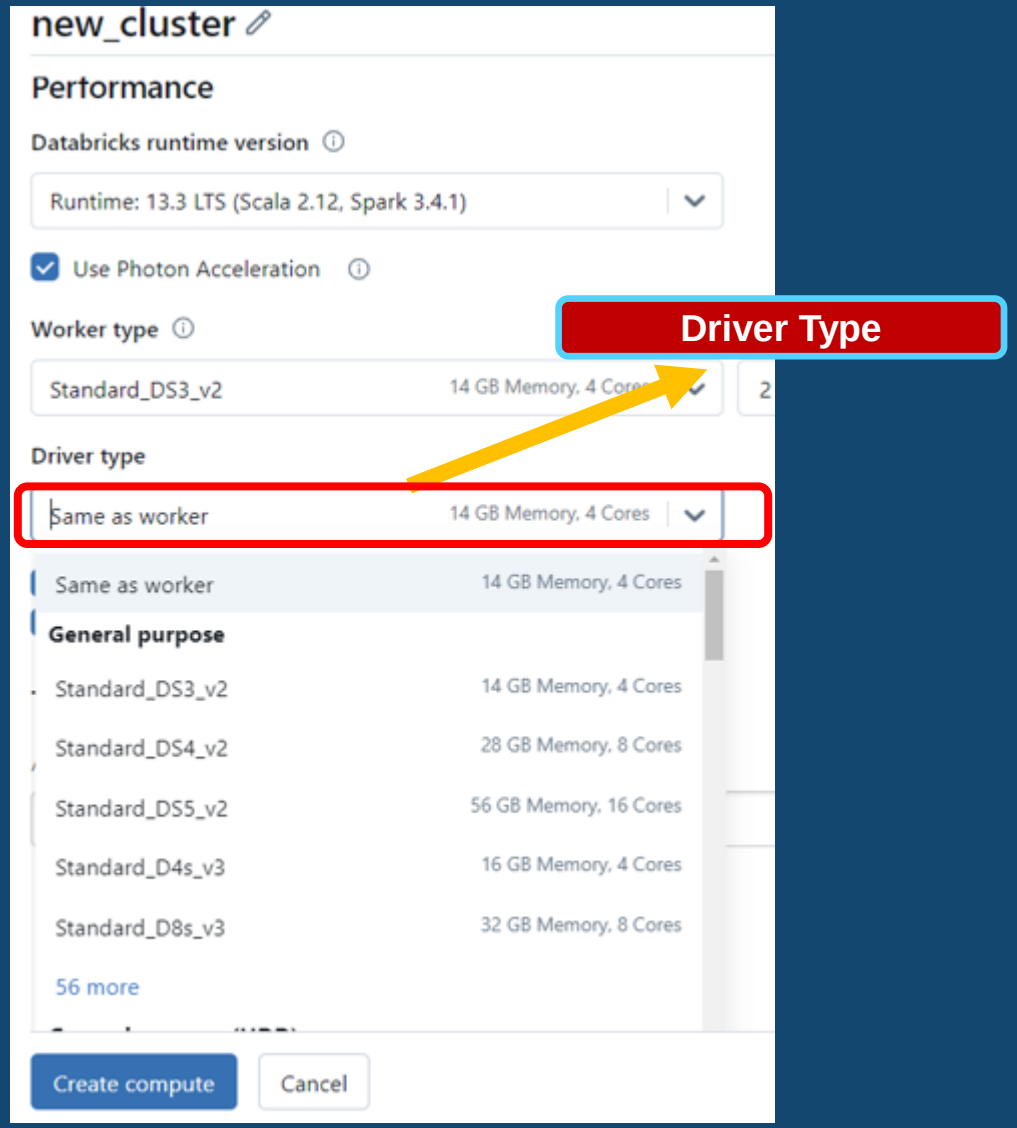

Driver Type

Make this larger if you want to do a lot of work on the driver, i.e., collecting data, local processing.

Cluster – Advanced Settings

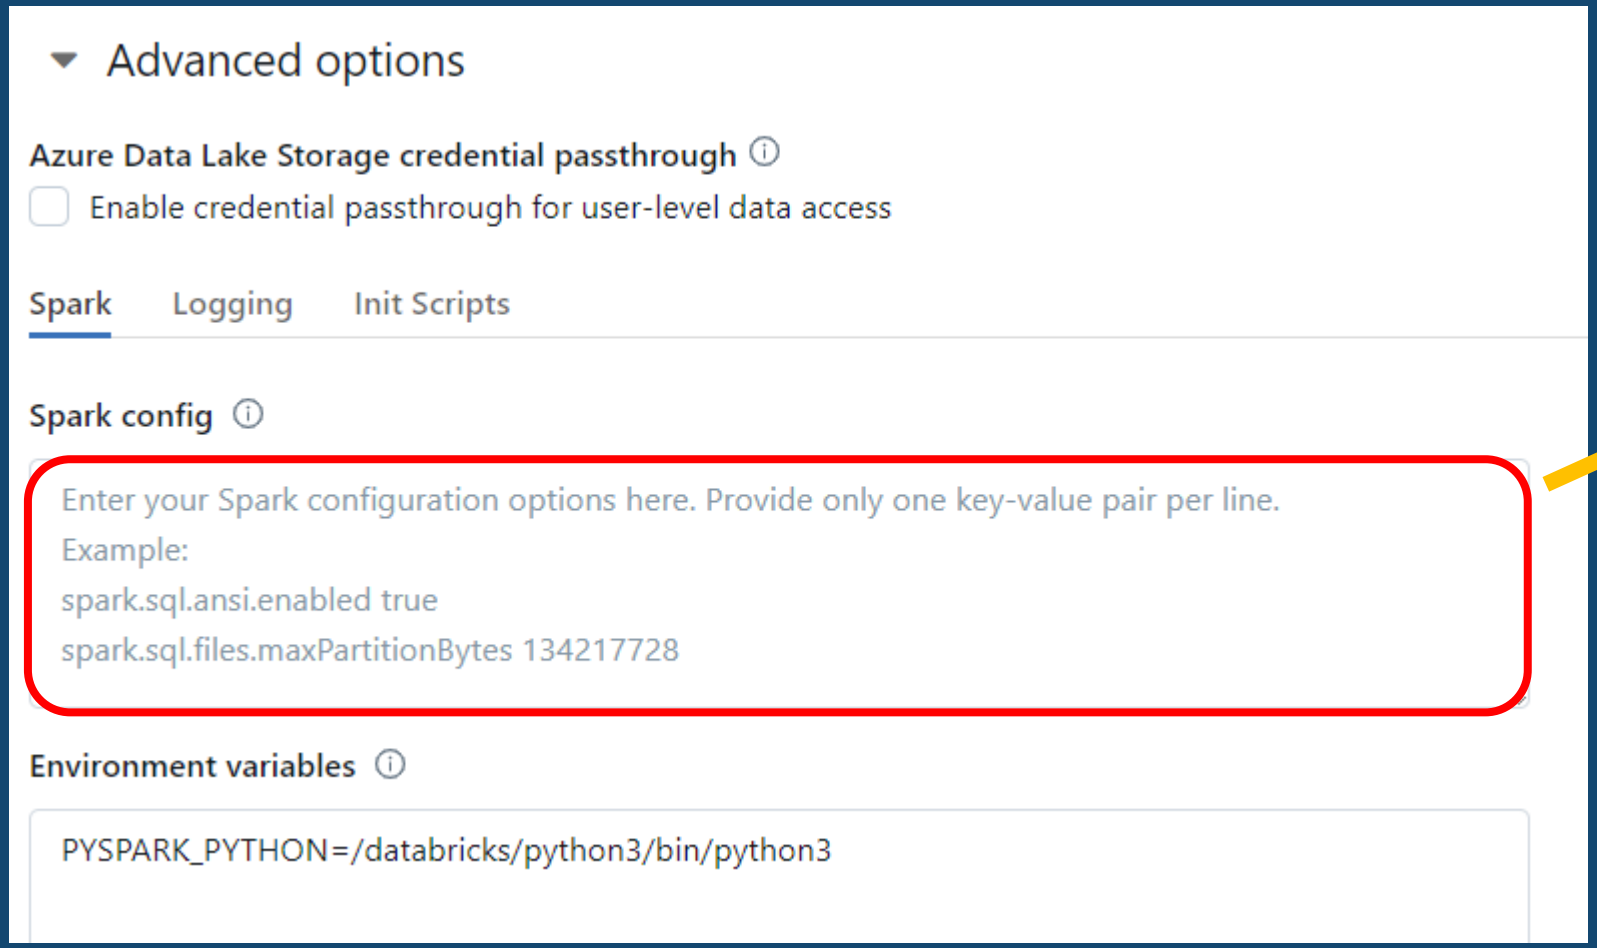

Modify Spark Configuration Settings to Improve performance

Cluster – Shuffles & Spills

Getting to the Spark UI from your Notebook

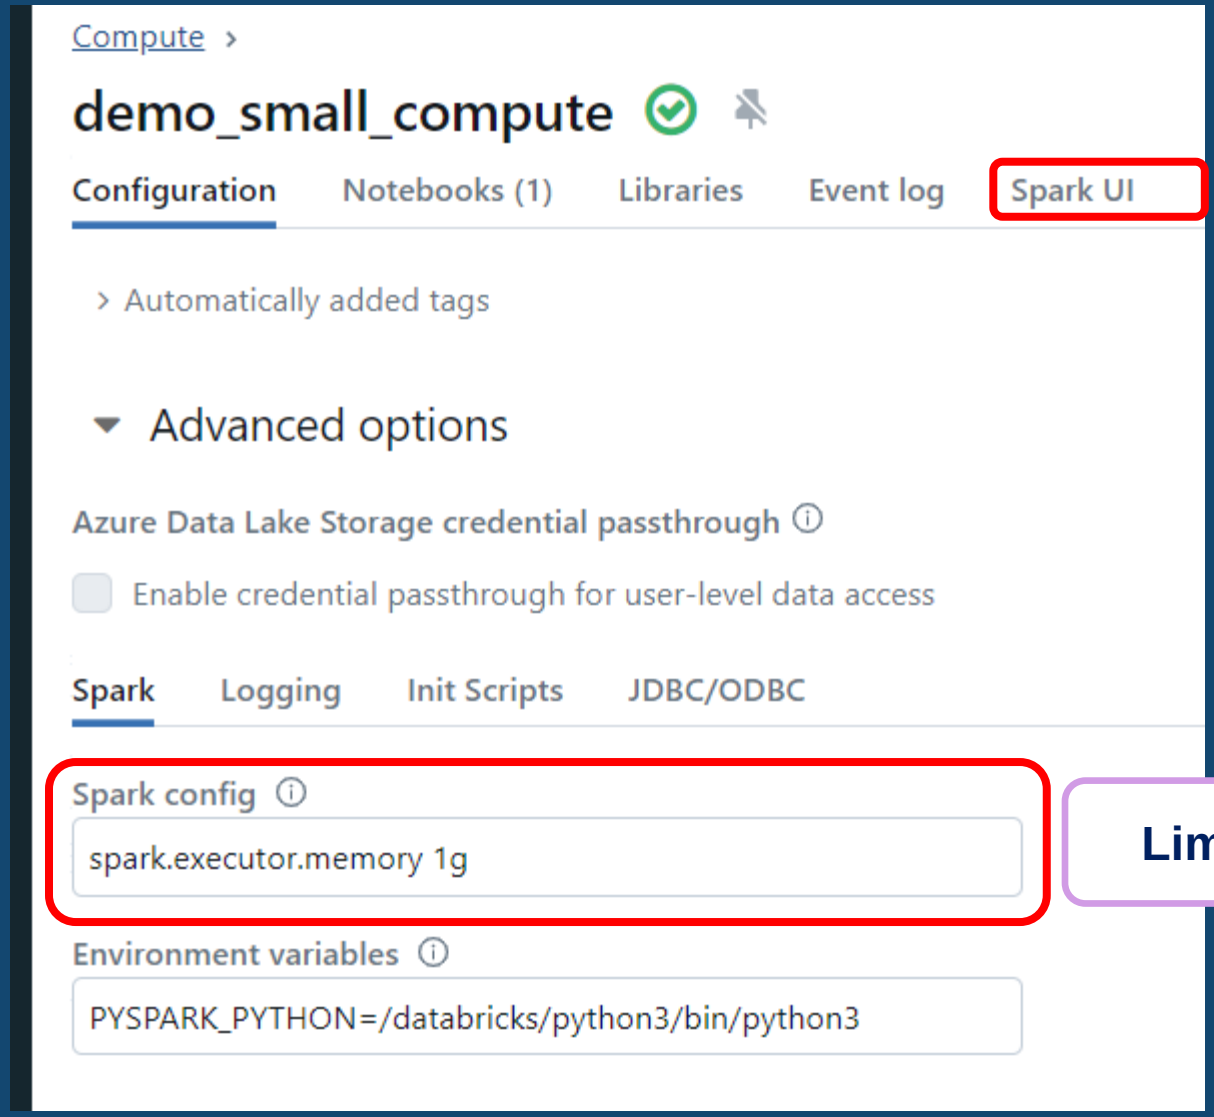

Click for the Spark UI

Limiting Memory to Force a Spill

Getting to the Spark UI from your Notebook

```
%python

spdf_sales = spark.sql('''
    SELECT o.o_orderstatus, c.c_name, count(*)
    FROM samples.tpch.lineitem       li
    left join samples.tpch.orders    o  ON (li.l_orderkey = o.o_orderkey)
    left join samples.tpch.customer  c  ON (o.o_custkey = c.c_custkey)
    GROUP BY o.o_orderstatus, c.c_name
    ''')
```

```
spdf_sales: pyspark.sql.dataframe.DataFrame
    o_orderstatus: string
    c_name: string
    count(1): long
```

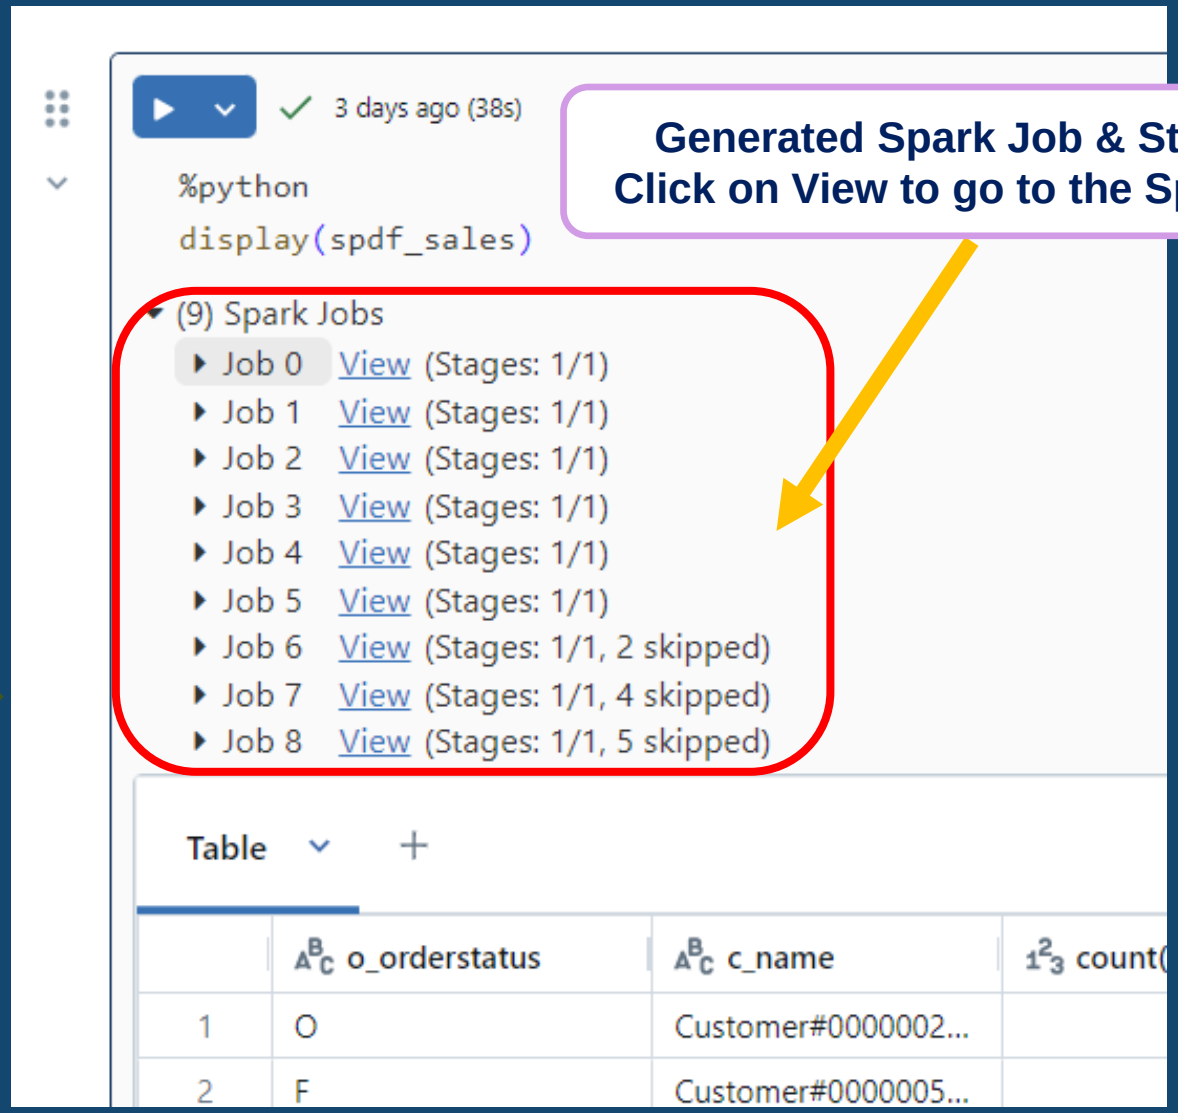

Generated Spark Job & Stages.
Click on View to go to the Spark UI.

Cluster – Shuffle & Spills

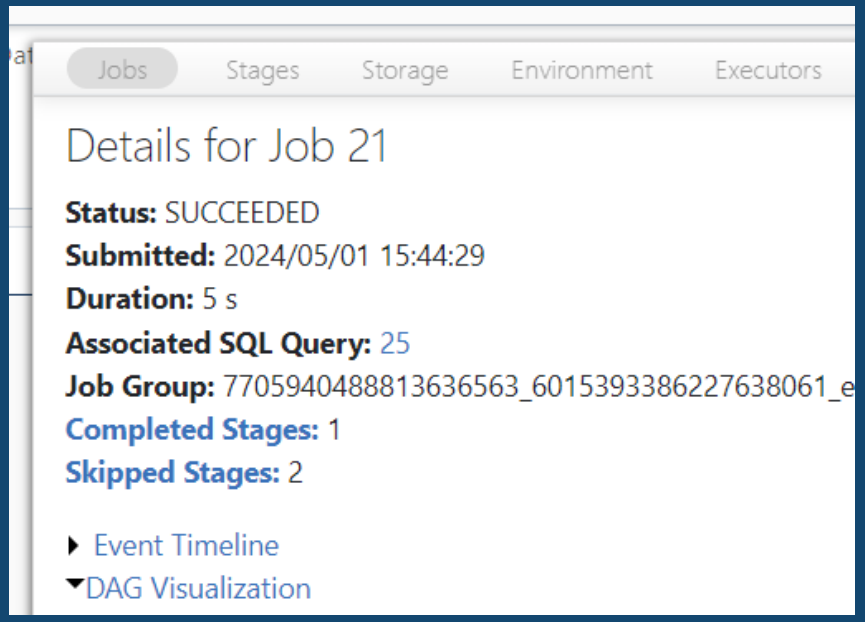

Shuffle Read/Writes

Completed Stages (1)

Page: 1 — 1 Pages. Jump to 1 . Show 100 items in a page. Go

Stage Id	Pool Name	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
40	7705940488813636563	display(spdf_sales) $anonfun$withThreadLocalCaptured$5 at LexicalThreadLocal.scala:63 +details	2024/05/01 15:44:29	5 s	9/9			150.2 MiB	81.8 MiB

Page: 1 — 1 Pages. Jump to 1 . Show 100 items in a page. Go

Cluster – Shuffle & Spills

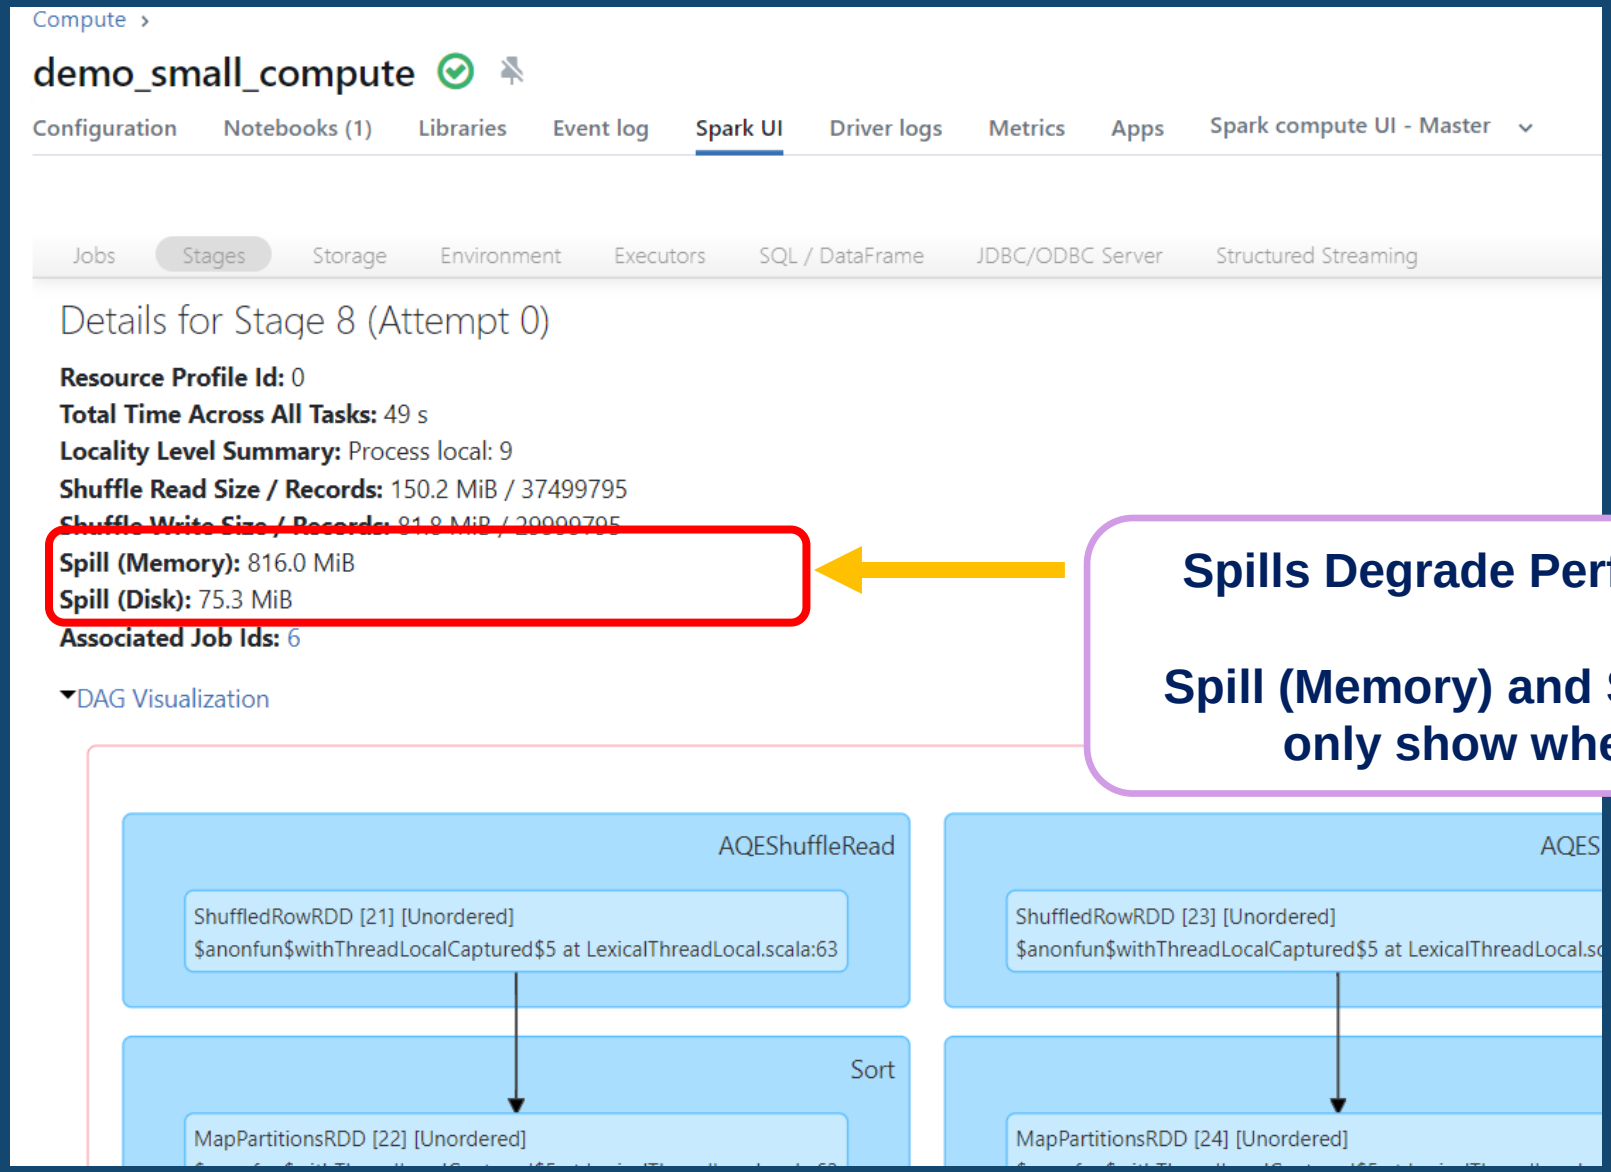

Spills Degrade Performance.

Spill (Memory) and Spill (Disk) only show when > 0

Wrapping Up

- Why Compute Resources Matter
- Data Standard or Premium
- Databricks/Apache Spark Cluster Architecture
- Hardware Under the Cluster Architecture
- Optimizing the Cluster Configuration Step By Step
- Shuffles & Spills

Thank You!